

Date: 9.15.2025

AWS Bedrock & Agentic AI

Contents

1. Introduction	2
2. Agentic AI in Bedrock - Amazon Bedrock AgentCore.....	2
3. Bedrock AgentCore Services.....	2
4. Bedrock AgentCore Pricing.....	3
5. Amazon Bedrock Agents.....	4
6. Strands Agent (SDK).....	5
7. Amazon Bedrock Agents vs Amazon Bedrock AgentCore.....	6
8. Conclusion	6

1. Introduction

AWS Bedrock: A secure, flexible service for building and deploying generative AI applications that provides access to top foundation models, tools for fine-tuning, safety, optimization, and integration with knowledge bases, enabling fast transition from testing to real-world use.

Agentic AI: Autonomous systems that plan, act, and adapt with minimal input. Using tools and data, they tackle complex tasks and boost efficiency—marking a shift as transformative as the internet or PC.

2. Agentic AI in Bedrock - Amazon Bedrock AgentCore

- Enables secure, scalable deployment of advanced AI agents
- Provides infrastructure tailored for dynamic workloads
- Offers tools to enhance agent capabilities
- Includes controls for real-world deployment
- Compatible with any framework (e.g., CrewAI, LangGraph, LlamaIndex, Strands Agents)
- Works with any foundation model, inside or outside Amazon Bedrock
- Removes the need to build specialized agent infrastructure
- Speeds up AI agent production

3. Bedrock AgentCore Services

Service	Description
Runtime	Secure, serverless environment for deploying and scaling dynamic AI agents and tools.
Browser	Cloud-based, secure browser runtime enabling agents to interact with websites at scale.
Code Interpreter	Allows agents to write and run code securely in sandboxed environments for complex tasks.
Gateway	Secure tool discovery, transforming APIs, Lambda functions, and services into agent-compatible tools easily.
Identity	Simplifies identity and access management, enabling secure access to AWS and third-party services.
Memory	Manages context-aware agent memory with full developer control, removing complex infrastructure needs.
Observability	Provides visibility for tracing, debugging, and monitoring agent workflows in production environments.

Here is the [Amazon Bedrock AgentCore \(Dev Guide\)](#) and [Amazon Bedrock AgentCore \(Overview\)](#)

4. Bedrock AgentCore Pricing

All services are priced on a consumption-based pricing scheme.

Service	Resource / Metric	Price
Runtime	CPU	\$0.0895 / ~₹7.90 per vCPU hour
	Memory	\$0.00945 / ~₹0.83 per GB hour
Browser Tool	CPU	\$0.0895 / ~₹7.90 per vCPU hour
	Memory	\$0.00945 / ~₹0.83 per GB hour
Code Interpreter	CPU	\$0.0895 / ~₹7.90 per vCPU hour
	Memory	\$0.00945 / ~₹0.83 per GB hour
Gateway	API Invocations (ListTools, InvokeTool)	\$0.005 / ~₹0.44 per 1,000 invocations
	Search API	\$0.025 / ~₹2.21 per 1,000 invocations
	Tool Indexing	\$0.02 / ~₹1.76 per 100 tools indexed per month
Identity	Token/API key requests (non-AWS resources)	\$0.010 / ~₹0.88 per 1,000 requests
		<i>Free via AgentCore Runtime or Gateway</i>
Memory	Short-Term Memory	\$0.25 / ~₹22.05 per 1,000 new events
	Long-Term Memory Storage (built-in strategy)	\$0.75 / ~₹66.16 per 1,000 memories/month
	Long-Term Memory Storage (custom strategy)	\$0.25 / ~₹22.05 per 1,000 memories/month
	Long-Term Memory Retrieval	\$0.50 / ~₹44.11 per 1,000 retrievals
Observability	Logs, metrics, traces	Charged per Amazon CloudWatch pricing

5. Amazon Bedrock Agents

Amazon Bedrock Agents is a fully managed service that builds agents using foundation models, APIs, and company data. It can:

- Break down complex user requests into steps
- Use retrieval-augmented generation (RAG) to pull in external info
- Retain memory for smoother multi-step interactions
- Automatically call APIs to take actions
- Support multi-agent collaboration for tackling complex tasks
- Provide built-in security and reliability via Guardrails

Feature	Description
Multi-Agent Collaboration	Multiple specialized agents work together under a supervisor to handle complex tasks.
Retrieval-Augmented Generation (RAG)	Agents access company data to enrich responses with accurate, real-time info.
Orchestrate and Execute	Agents break down tasks, call APIs, and automate workflows using FM reasoning.
Memory Retention	Agents retain context across interactions for personalized, seamless experiences.
Code Interpretation	Agents generate and execute code securely for analytics, math, and data tasks.

[AI Agents – Amazon Bedrock Agents – AWS](#)

6. Strands Agent (SDK)

Strands Agents: A lightweight, model-driven SDK for building AI agents with minimal code that supports everything from simple assistants to complex autonomous workflows and scales from local development to production.

Feature	Description
Model Agnostic	Supports Amazon Bedrock, Anthropic, LiteLLM, Llama, Ollama, OpenAI, Writer, and custom providers
Flexible & Lightweight	Simple, customizable agent loop
Multi-Agent & Streaming	Build autonomous, collaborative agents with streaming support
Built-in MCP Support	Integrates Model Context Protocol servers with thousands of tools
Python-Based Tools	Easily create tools using Python decorators
Hot Reloading	Auto-reload tools from the ./tools/ directory
Multiple Model Providers	Includes Amazon Bedrock, Ollama, LlamaAPI, and more
Pre-built Tools	Optional strands-agents-tools package for quick experimentation

[Strands Agents Docs](#)

7. Amazon Bedrock Agents vs Amazon Bedrock AgentCore

Feature	Amazon Bedrock Agents	Amazon Bedrock AgentCore
Interface	GUI-based setup via AWS Console, minimal coding required.	Web-based console, extensive coding and infrastructure management required.
Agent Creation	Build agents using templates and natural language instructions.	Requires defining agent logic, containerization, and deployment using SDK, CLI, or APIs.
Customization	Limited customization, suitable for standard workflows.	Full control over agent behavior and integration with external frameworks.
Tool Integration	Limited, primarily through pre-configured connectors.	Extensive via supports custom APIs, Lambda functions, and third-party services.
Memory Management	Basic memory retention for context-aware interactions.	Advanced memory management with short-term and long-term storage options.
Code Execution	Supports dynamic code generation and execution in a secure environment.	Includes a Code Interpreter tool for executing code in isolated sandbox environments.
Browser Interaction	Limited browser interaction capabilities.	Includes a secure, cloud-based browser runtime for web-based workflows.
Observability	Basic monitoring through AWS CloudWatch.	Comprehensive observability with detailed dashboards and telemetry support.
Deployment	Simplified deployment process via AWS Console.	Requires manual setup and configuration for deployment and scaling.

8. Conclusion

- **Strands Agents SDK:** Ideal for developing and prototyping AI agents locally with minimal code and high flexibility.
- **Amazon Bedrock Agent Builder:** A no-code/low-code visual tool to quickly create and customize AI agents via a GUI, enabling faster development without deep coding.
- **Amazon Bedrock AgentCore:** Built for deploying, scaling, and managing AI agents securely and reliably in production environments.