

Date: 04.07.2025

# K8s Service Management and Debugging Essentials

## Contents

1. Service Types .....	2
a. ClusterIP .....	2
b. NodePort .....	2
c. LoadBalancer .....	2
d. ExternalName .....	2
2. Namespaces .....	3
a. What are namespaces? .....	3
b. Initial Namespaces .....	3
c. Manifest file to create a namespace .....	3
3. Debugging running pods .....	4
4. Resource Cleanup .....	4

## 1. Service Types

```
apiVersion: v1
kind: Service
metadata:
  name: <SERVICE_NAME>
spec:
  selector:
    app: <DEPLOYMENT.METADATA.LABELS.APP>
  type: <SERVICE_TYPE>
  ports:
    - protocol: <PROTOCOL>
      port: <PORT_WHERE_SERVICE_EXPOSED_IN_CLUSTER>
      targetPort: <PORT_ON_POD_THAT_RECIEVES_TRAFFIC>
```

### a. ClusterIP

- Exposes the Service on a cluster-internal IP.
- This value makes the Service only reachable from within the cluster.
- This is the default that is used if you don't explicitly specify a type for a Service.
- You can expose the Service to the public internet using an Ingress or a Gateway.
- Assigns an IP address from a pool of IP addresses that your cluster has reserved for that purpose.

### b. NodePort

- Exposes the Service on each Node's IP at a static port (the NodePort).
- To make the node port available, Kubernetes sets up a cluster IP address, the same as if you had requested a Service of type: ClusterIP.
- The Kubernetes control plane allocates a port from a range specified by --service-node-port-range flag (default: 30000-32767).

### c. LoadBalancer

- Exposes the Service externally using an external load balancer.
- Kubernetes does not directly offer a load balancing component; you must provide one, or you can integrate your Kubernetes cluster with a cloud provider.
- On cloud providers which support external load balancers, setting the `type` field to LoadBalancer provisions a load balancer for your Service.

### d. ExternalName

- Maps the Service to the contents of the externalName field (for example, to the hostname api.foo.bar.example).
- The mapping configures your cluster's DNS server to return a CNAME record with that external hostname value. No proxying of any kind is set up.
- Accepts an IPv4 address string, but treats that string as a DNS name comprised of digits.

## 2. Namespaces

### a. What are namespaces?

- *namespaces* provide a mechanism for isolating groups of resources within a single cluster.
- Names of resources need to be unique within a namespace, but not across namespaces.
- Namespace-based scoping is applicable only for namespaced objects (*e.g. Deployments, Services, etc.*) and not for cluster-wide objects (*e.g. StorageClass, Nodes, PersistentVolumes, etc.*).
- It is intended for use in environments with many users spread across multiple teams, or projects.
- It is not necessary to use multiple namespaces to separate slightly different resources, such as different versions of the same software: use `labels` to distinguish resources within the same namespace.
- It is a way to divide cluster resources between multiple users (via resource quota).

### b. Initial Namespaces

Namespace	Summary
default	Used for resources when no other namespace is specified.
kube-node-lease	Stores node heartbeat data to help detect node failures.
kube-public	Publicly readable namespace, typically for cluster-wide shared resources.
kube-system	Contains system components and resources created by Kubernetes itself.

### c. Manifest file to create a namespace

```
apiVersion: v1
kind: Namespace
metadata:
  name: <NAME>
  labels:
    name: <NAME>
```

### 3. Debugging running pods

`kubectl get pods` → Get status of all pods

`kubectl describe pod <PODNAME>` → More information like configuration & status info  
Analyze 'Event' section for any debug information

`kubectl get events` → Lists all the events

`kubectl get pod <PODNAME> -o yaml` → Returns in YAML

`kubectl logs <PODNAME> <CONTAINERNAME>` → View logs of affected container

`kubectl logs --previous <PODNAME> -c <CONTAINERNAME>` → Previous container's logs

`kubectl exec <PODNAME> -c <CONTAINERNAME> -- <CMD>` → Run command inside a specific container

`kubectl exec -it <PODNAME> -- /bin/bash` → Open an interactive terminal connecting to the pod

`kubectl exec --stdin --tty <PODNAME> -- /bin/bash` → Same as previous command

`kubectl debug node/<NODENAME> -it --image=ubuntu` → Interactive shell to the node  
Find the node that is running the pod first

### 4. Resource Cleanup

`kubectl delete <RESOURCE> <NAME>` → Delete specific resource

`kubectl get pods --field-selector=status.phase==<Succeeded/Failed>` → Get completed/failed pods

`kubectl delete pod <PODNAME>` → Use the result of above command and then delete pod

`kubectl apply --prune -l app=myapp -f ./manifests --all` → Delete any previous resources in the cluster with label [app=myapp] but not defined in the current set of manifests being applied

`kubectl delete ns <ns-name>` → Delete entire namespace and all its resources

```

spec:

ttlSecondsAfterFinished: <TIME\_IN\_SECS>

``` → In Jobs and CronJobs, delete that resource after it is finished